

## DWDM PROJECT

### *Importing required modules*

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from xgboost import XGBRegressor
from sklearn.ensemble import VotingRegressor, RandomForestRegressor,
GradientBoostingRegressor, StackingRegressor
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import mean_squared_error,
r2_score, accuracy_score, classification_report
from sklearn.preprocessing import StandardScaler, LabelEncoder
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
import warnings
warnings.filterwarnings("ignore")
```

### *A small subset of data-set*

```
df = pd.read_csv(r"D:\3-1\Theory\F2 DWDM\project\data.csv")
df.head(10)
```

	mandal	date	rainfall	temp_min	temp_max	humidity_min	\
0	Shaikpet	01-01-2018	0.0	17.4	30.30	40.6	
1	Shaikpet	02-01-2018	0.0	19.4	31.60	37.5	
2	Shaikpet	03-01-2018	0.0	19.3	29.90	43.8	
3	Shaikpet	04-01-2018	0.0	18.2	29.30	41.6	
4	Shaikpet	05-01-2018	0.0	17.1	28.00	42.9	
5	Shaikpet	06-01-2018	0.0	15.9	29.30	37.5	
6	Shaikpet	07-01-2018	0.0	17.2	29.90	35.5	
7	Shaikpet	08-01-2018	0.0	16.9	27.89	40.9	
8	Shaikpet	09-01-2018	0.0	17.3	26.70	42.7	
9	Shaikpet	10-01-2018	0.0	17.4	26.40	51.2	

	humidity_max	wind_speed_min	wind_speed_max
0	85.0	0.0	4.8
1	74.9	0.0	8.5
2	86.4	0.0	8.0
3	80.9	0.0	7.0
4	71.2	0.0	11.2
5	76.8	0.0	7.3
6	74.4	0.0	7.1
7	74.6	0.0	10.3
8	81.6	0.0	9.8
9	77.8	0.0	13.2

### Statistics of the data-set

```
df.info()
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 35042 entries, 0 to 35041
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   mandal                35042 non-null  object
1   date                  35042 non-null  object
2   rainfall              35042 non-null  float64
3   temp_min              35042 non-null  float64
4   temp_max              35042 non-null  float64
5   humidity_min          35042 non-null  float64
6   humidity_max          35042 non-null  float64
7   wind_speed_min        34001 non-null  float64
8   wind_speed_max        34977 non-null  float64
dtypes: float64(7), object(2)
memory usage: 2.4+ MB
```

	rainfall	temp_min	temp_max	humidity_min
count	35042.000000	35042.000000	35042.000000	35042.000000
mean	2.736562	23.189849	33.593165	42.558944
std	9.082997	3.474841	3.941404	18.773714
min	0.000000	6.000000	19.500000	0.000000
25%	0.000000	21.500000	30.900000	27.300000
50%	0.000000	23.500000	33.000000	42.000000
75%	0.100000	25.200000	36.400000	56.000000
max	149.700000	33.300000	44.800000	99.300000

	wind_speed_min	wind_speed_max
count	34001.000000	34977.000000
mean	0.385418	10.433857
std	1.056075	11.976747
min	0.000000	0.000000
25%	0.000000	4.400000
50%	0.000000	7.900000
75%	0.100000	11.800000
max	15.500000	234.400000

### *Removing duplicate records*

AS there are no null values we are skipping the step which fills the null values

```
df = df.drop_duplicates()
```

### *Counting null values of each attribute in the data-set*

```
df.isnull().sum()
mandal      0
date        0
rainfall    0
temp_min    0
temp_max    0
humidity_min 0
humidity_max 0
wind_speed_min 1041
wind_speed_max 65
dtype: int64
```

### *Replacing the null values with the mean*

```
df["wind_speed_max"] = df["wind_speed_max"].fillna(df["wind_speed_max"].mean())
df["wind_speed_min"] = df["wind_speed_min"].fillna(df["wind_speed_min"].mean())
# df = df.drop(columns=['wind_speed_min'])
```

### *Statistics of data-set*

```
df.describe()
```

	rainfall	temp_min	temp_max	humidity_min
count	35042.000000	35042.000000	35042.000000	35042.000000
mean	2.736562	23.189849	33.593165	42.558944
std	9.082997	3.474841	3.941404	18.773714
min	0.000000	6.000000	19.500000	0.000000
25%	0.000000	21.500000	30.900000	27.300000
50%	0.000000	23.500000	33.000000	42.000000
75%	0.100000	25.200000	36.400000	56.000000
max	80.437477	35.000000	39.000000	92.400000


max	149.700000	33.300000	44.800000	99.300000
100.000000				


	wind_speed_min	wind_speed_max
count	35042.000000	35042.000000
mean	0.385418	10.433857
std	1.040270	11.965633
min	0.000000	0.000000
25%	0.000000	4.400000
50%	0.000000	7.900000
75%	0.200000	11.800000
max	15.500000	234.400000

*Box plot of temp\_min vs temp\_max*

```
plt.figure(figsize=(6, 6))
sns.boxplot(data=df[['temp_min', 'temp_max']])
plt.title("Box plots between tempearture")
plt.show()
plt.figure(figsize=(6, 6))
sns.boxplot(data=df[['humidity_min', 'humidity_max']])
plt.title("Box plots between tempearture")
plt.show()
```

Box plots between tempearture





*Removing the outliers*

```
# def winsorize_outliers(df, columns, lower_quantile=0.05,
upper_quantile=0.95):
#     for column in columns:
#         lower_bound = df[column].quantile(lower_quantile)
#         upper_bound = df[column].quantile(upper_quantile)
#         df[column] = np.where(df[column] < lower_bound, lower_bound,
df[column])
#         df[column] = np.where(df[column] > upper_bound, upper_bound,
df[column])
#     return df

# columns_to_clean = ["rainfall", "temp_min", "temp_max",
"humidity_min", "humidity_max", "wind_speed_min", "wind_speed_max"]
# df= winsorize_outliers(df, columns_to_clean)
```


```
def iqr_outlier_handling(df, columns):
    for column in columns:
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        df[column] = np.where(df[column] < lower_bound, lower_bound,
df[column])
        df[column] = np.where(df[column] > upper_bound, upper_bound,
df[column])
    return df

columns_to_clean = ["rainfall", "temp_min", "temp_max",
"humidity_min", "humidity_max", "wind_speed_min", "wind_speed_max"]
df = iqr_outlier_handling(df, columns_to_clean)
```

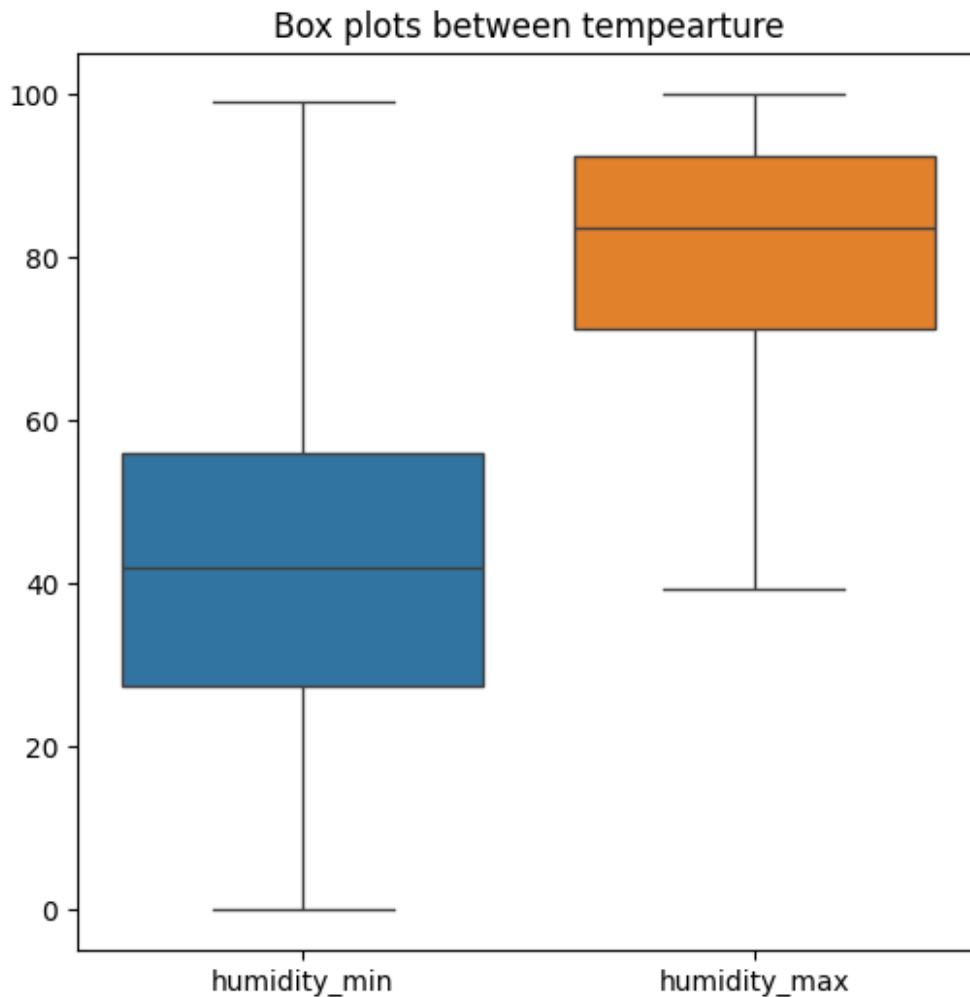
*Box plot after the removal of outliers*

```
plt.figure(figsize=(6, 6))
sns.boxplot(data=df[['temp_min', 'temp_max']])
plt.title("Box plots between tempearture")
plt.show()
plt.figure(figsize=(6, 6))
sns.boxplot(data=df[['humidity_min', 'humidity_max']])
plt.title("Box plots between tempearture")
plt.show()
```

Box plots between tempearture







*Predicting the type of weather*

```
def categorize_weather(row):  
    rainfall = row['rainfall']  
    humidity_min = row['humidity_min']  
    wind_speed_max = row['wind_speed_max']  
  
    if rainfall == 0 and humidity_min < 50:  
        return 'Sunny'  
    elif rainfall == 0 and 50 <= humidity_min < 70:  
        return 'Partly Sunny'  
    elif rainfall == 0 and humidity_min >= 70:  
        return 'Partly Cloudy'  
    elif 0 < rainfall <= 10 and humidity_min > 60:  
        return 'Sun and Rain'  
    elif rainfall > 10 and humidity_min > 70:  
        return 'Heavy Raining'  
    elif 1 <= rainfall <= 10 and humidity_min > 60:  
        return 'Light Raining'
```

```

elif rainfall > 10 and humidity_min > 80 and wind_speed_max > 15:
    return 'Thunderstorms'
elif rainfall == 0 and humidity_min > 60:
    return 'Cloudy'
elif wind_speed_max > 15:
    return 'Windy'
elif rainfall <= 1 and 50 <= humidity_min < 70:
    return 'Rainbow'
else:
    return 'Clear'

```

```
df['weather_type'] = df.apply(categorize_weather, axis=1)
```

```
print(df[['rainfall', 'humidity_min', 'temp_min', 'wind_speed_max',
'weather_type']].head())
```

	rainfall	humidity_min	temp_min	wind_speed_max	weather_type
0	0.0	40.6	17.4	4.8	Sunny
1	0.0	37.5	19.4	8.5	Sunny
2	0.0	43.8	19.3	8.0	Sunny
3	0.0	41.6	18.2	7.0	Sunny
4	0.0	42.9	17.1	11.2	Sunny

*Summarization of the data-set*

```
df.describe()
```

	rainfall	temp_min	temp_max	humidity_min
humidity_max \				
count	35042.000000	35042.000000	35042.000000	35042.000000
mean	0.063146	23.274716	33.594616	42.558933
std	0.107586	3.189836	3.936952	18.773680
min	0.000000	15.950000	22.650000	0.000000
25%	0.000000	21.500000	30.900000	27.300000
50%	0.000000	23.500000	33.000000	42.000000
75%	0.100000	25.200000	36.400000	56.000000
max	0.250000	30.750000	44.650000	99.050000
	wind_speed_min	wind_speed_max		
count	35042.000000	35042.000000		

mean	0.118944	8.746330
std	0.196009	5.735108
min	0.000000	0.000000
25%	0.000000	4.400000
50%	0.000000	7.900000
75%	0.200000	11.800000
max	0.500000	22.900000

*Finding the Accuracy of Rainfall using Stacking, Random forest, XG Boost and GradientBoost Models*

RAINFALL(Stacking,Random forest,XGboost,GradientBoost)

```
X = df.drop(columns=['rainfall', 'mandal', 'date', 'weather_type'])
y = df['rainfall']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.25, random_state=42)

rf = RandomForestRegressor(n_estimators=100, random_state=42)
xgb = XGBRegressor(n_estimators=100, random_state=42)
gb = GradientBoostingRegressor(n_estimators=100, random_state=42)

estimators = [
    ('rf', rf),
    ('xgb', xgb),
    ('gb', gb)
]

stacking_regressor = StackingRegressor(estimators=estimators,
final_estimator=LinearRegression())
stacking_regressor.fit(X_train, y_train)

y_pred_stacking = stacking_regressor.predict(X_test)

rmse_stacking = np.sqrt(mean_squared_error(y_test, y_pred_stacking))
accuracy_stacking = r2_score(y_test, y_pred_stacking) * 100

print(f'Accuracy (R^2 Score) for Stacking rainfall prediction:
{accuracy_stacking:.2f}%')

Accuracy (R^2 Score) for Stacking rainfall prediction: 50.33%
```

*Finding the Accuracy of temp\_max using Random forest Model*

temp\_max(Random forest)

```

X = df.drop(columns=['temp_max', 'mandal', 'date', 'weather_type'])
y = df['temp_max']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.2, random_state=42)

model = RandomForestRegressor(n_estimators=100, random_state=42,
max_depth=15)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

rmse = np.sqrt(mean_squared_error(y_test, y_pred))
accuracy_temp_max = r2_score(y_test, y_pred) * 100

print(f'Accuracy (R^2 Score) for temp_max prediction:
{accuracy_temp_max:.2f}%')

Accuracy (R^2 Score) for temp_max prediction: 87.29%

```

*Finding the Accuracy of temp\_min using Random forest Model*

temp\_min(Random forest)

```

X = df.drop(columns=['temp_min', 'mandal', 'date', 'weather_type'])
y = df['temp_min']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.2, random_state=42)

model = RandomForestRegressor(n_estimators=100, random_state=42,
max_depth=15)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

rmse = np.sqrt(mean_squared_error(y_test, y_pred))
accuracy_temp_min = r2_score(y_test, y_pred) * 100

print(f'Accuracy (R^2 Score) for temp_min prediction:
{accuracy_temp_min:.2f}%')

Accuracy (R^2 Score) for temp_min prediction: 79.54%

```

*Finding the Accuracy of humidity\_min using GradientBoostingRegressor Model*

humidity\_min (GradientBoostingRegressor)

```
X = df.drop(columns=['humidity_min', 'mandal', 'date',
'weather_type'])
y = df['humidity_min']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.2, random_state=42)

model = GradientBoostingRegressor(n_estimators=100, learning_rate=0.1,
max_depth=3, random_state=42)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

rmse = np.sqrt(mean_squared_error(y_test, y_pred))
accuracy_humidity_min = r2_score(y_test, y_pred) * 100

print(f'Accuracy (R^2 Score) for humidity_min prediction:
{accuracy_humidity_min:.2f}%')

Accuracy (R^2 Score) for humidity_min prediction: 87.70%
```

*Finding the Accuracy of humidity\_max using GradientBoostingRegressor Model*

```
X = df.drop(columns=['humidity_max', 'mandal', 'date',
'weather_type'])
y = df['humidity_max']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.25, random_state=42)

model = RandomForestRegressor(n_estimators=100, random_state=42,
max_depth=15)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

rmse = np.sqrt(mean_squared_error(y_test, y_pred))
accuracy_humidity_max = r2_score(y_test, y_pred) * 100

print(f'Accuracy (R^2 Score) for humidity_max prediction:
{accuracy_humidity_max:.2f}%')

Accuracy (R^2 Score) for humidity_max prediction: 71.46%
```

*Finding the Accuracy of wind\_speed\_max using Voting, Random forest, XGboost and GradientBoost Models*

Wind\_speed\_max(Voting,Random forest,XGboost,GradientBoost)

```
X = df.drop(columns=['rainfall', 'mandal', 'date', 'weather_type'])
y = df['wind_speed_max']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.25, random_state=42)

rf = RandomForestRegressor(n_estimators=100, random_state=42,
max_depth=15)
xgb = XGBRegressor(n_estimators=100, random_state=42, max_depth=5)
gb = GradientBoostingRegressor(n_estimators=100, random_state=42,
max_depth=5)

voting_regressor = VotingRegressor(estimators=[('rf', rf), ('xgb',
xgb), ('gb', gb)])
voting_regressor.fit(X_train, y_train)

y_pred = voting_regressor.predict(X_test)

rmse = np.sqrt(mean_squared_error(y_test, y_pred))
accuracy_wind_max = r2_score(y_test, y_pred) * 100

print(f'Accuracy (R^2 Score) for predicting wind_speed_max:
{accuracy_wind_max:.2f}%')

Accuracy (R^2 Score) for predicting wind_speed_max: 100.00%
```

*Finding the Accuracy of wind\_speed\_min using Voting, Random forest, XGboost and GradientBoost Models*

Wind\_speed\_min(Voting,Random forest,XGboost,GradientBoost)

```
X = df.drop(columns=['rainfall', 'mandal', 'date', 'weather_type'])
y = df['wind_speed_min']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.25, random_state=42)

rf = RandomForestRegressor(n_estimators=100, random_state=42,
max_depth=15)
xgb = XGBRegressor(n_estimators=100, random_state=42, max_depth=5)
```

```

gb = GradientBoostingRegressor(n_estimators=100, random_state=42,
max_depth=5)

voting_regressor = VotingRegressor(estimators=[('rf', rf), ('xgb',
xgb), ('gb', gb)])
voting_regressor.fit(X_train, y_train)

y_pred = voting_regressor.predict(X_test)

rmse = np.sqrt(mean_squared_error(y_test, y_pred))
accuracy_wind_min = r2_score(y_test, y_pred) * 100

print(f'Accuracy (R^2 Score) for predicting wind_speed_min:
{accuracy_wind_min:.2f}%')

Accuracy (R^2 Score) for predicting wind_speed_min: 100.00%

```

*Finding the Accuracy of weather\_type using DecisionTreeClassifier Model*

```

X = df.drop(columns=['rainfall', 'mandal', 'date', 'weather_type'])
y = df['weather_type']

label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled,
y_encoded, test_size=0.25, random_state=42)

dt = DecisionTreeClassifier(random_state=42, max_depth=5)

dt.fit(X_train, y_train)

y_pred = dt.predict(X_test)

accuracy = accuracy_score(y_test, y_pred) * 100
report = classification_report(y_test, y_pred,
target_names=label_encoder.classes_)

print(f'Accuracy for predicting weather_type: {accuracy:.2f}%')

Accuracy for predicting weather_type: 84.37%

```

*Histogram representation of each and every attribute*

```

columns_to_plot = ['rainfall', 'temp_min', 'temp_max', 'humidity_min',
'humidity_max', 'wind_speed_min', 'wind_speed_max']

```




```
colors = ['blue', 'orange', 'green', 'red', 'purple', 'brown', 'pink']

n = len(columns_to_plot)
plt.figure(figsize=(6, 4 * n))

for i, (column, color) in enumerate(zip(columns_to_plot, colors),
start=1):
    plt.subplot(n, 1, i)
    plt.hist(df[column], bins=30, color=color, alpha=0.7,
edgecolor='black')
    plt.title(f'Histogram of {column}')
    plt.xlabel(column)
    plt.ylabel('Frequency')
    # plt.grid(axis='y', alpha=0.75)

plt.tight_layout()
plt.show()
```






### Heatmap of Correlation Matrix After Removing Outliers

```
numeric_df = df.select_dtypes(include=['int64', 'float64'])
correlation_matrix = numeric_df.corr()

plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
            fmt=".2f", square=True)
plt.title('Heatmap of Correlation Matrix After Removing Outliers')
plt.show()
```



### Representation of each and every attribute's Accuracy and the model used for it


```
accuracies = {
    'Rainfall Prediction (Stacking)': accuracy_stacking,
```

```

'Temp Max Prediction (RandomForest)': accuracy_temp_max,
'Temp Min Prediction (RandomForest)': accuracy_temp_min,
'Humidity Min Prediction (GradientBoosting)':
accuracy_humidity_min,
'Humidity Max Prediction (RandomForest)': accuracy_humidity_max,
'Wind Speed Max Prediction (Voting)': accuracy_wind_max,
'Wind Speed Min Prediction (Voting)': accuracy_wind_min,
'Weather Type Prediction (DecisionTree)': accuracy
}

fig, ax = plt.subplots(figsize=(10, 6))
ax.bar(accuracies.keys(), accuracies.values(), color='skyblue')
ax.set_title('Model Accuracies')
ax.set_xlabel('Prediction Models')
ax.set_ylabel('Accuracy (%)')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()

```



*Plot of all attribute's accuracies together*

```

model_names = [
    'Rainfall',
    'Temp Max',
    'Temp Min',

```

```

        'Humidity Min ',
        'Humidity Max ',
        'Wind Speed Max ',
        'Wind Speed Min ',
        'Weather Type'
    ]
    accuracies = [
        accuracy_stacking,
        accuracy_temp_max,
        accuracy_temp_min,
        accuracy_humidity_min,
        accuracy_humidity_max,
        accuracy_wind_max,
        accuracy_wind_min,
        accuracy
    ]

plt.figure(figsize=(12, 6))
plt.plot(model_names, accuracies, marker='o', color='b',
linestyle='-', linewidth=2, markersize=6)

for i, (model, acc) in enumerate(zip(model_names, accuracies)):
    plt.text(i, acc + 1, f'{model}\n{acc:.2f}%', ha='center',
va='bottom', fontsize=9)
    plt.plot(i, acc, marker='o', markersize=10,
markerfacecolor='none', markeredgecolor='red')

plt.title('Model Accuracies')
plt.xlabel('Prediction Models')
plt.ylabel('Accuracy (%)')
plt.xticks(rotation=45, ha='right')
plt.ylim(0, 110)
plt.grid(True)
plt.tight_layout()
plt.show()

```

